

Transfer a photo into a Monet style picture

(This part we got inspiration and learning from <https://www.kaggle.com/basu369victor/style-transfer-deep-learning-algorithm> (<https://www.kaggle.com/basu369victor/style-transfer-deep-learning-algorithm>))

First, we have a original image and a style image . Then, we still use VGG16 but only pick the convolution layers for extracting image's features and delete the layers for classification. Finally, we will get a new picture combine the style of the style image and the objects from original image .

```
In [1]: import os
import numpy as np
import tensorflow
from tensorflow.keras import applications
from tensorflow.keras.applications import vgg16
from tensorflow.keras.models import Model
from scipy.optimize import fmin_l_bfgs_b
from tensorflow.keras.preprocessing.image import load_img, save_img, img_to_array
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: import tensorflow.keras.backend as K
K.clear_session()
ori_path = './photo.jpg'
style_path = './Claude_Monet.jpg'
```

Generation image's size

```
In [3]: ori_image = load_img(ori_path)
nrow, ncol = ori_image.size
nrow = nrow // 10
ncol = ncol // 10
print(str(nrow), str(ncol))
```

403 302

Show 2 images

```
In [4]: ori_image = load_img(ori_path,target_size=(ncol, nrow))
print(ori_image.size)
style_image = load_img(style_path, target_size=(ncol, nrow))
plt.axis('off')
plt.title('Photo')
plt.imshow(ori_image)
plt.show()
plt.axis('off')
plt.title('Claude_Monet')
plt.imshow(style_image)
```

(403, 302)

Photo



Out[4]: <matplotlib.image.AxesImage at 0x7f028571a358>

Claude_Monet



Pre-Process the images as the VGG16 input.Using tensor to represent the image,and bulid a new imgae as ouput

```
In [5]: def preprocess(image):  
        img = img_to_array(image)  
        img = vgg16.preprocess_input(img[np.newaxis,:])  
        img = K.variable(img)  
        return img  
ori_img = preprocess(ori_image)  
style_img = preprocess(style_image)  
result_image = K.placeholder((1,ncol, nrow,3)) # 'channels_last' format()  
result_image.shape
```

WARNING:tensorflow:From /usr/local/lib/python3.5/dist-packages/tensorflow/python/ops/resource_variable_ops.py:435: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

```
Out[5]: TensorShape([Dimension(1), Dimension(302), Dimension(403), Dimension(3)])
```

Then intergrating original image, style image and result image.

```
In [6]: vgg16_input = K.concatenate([ori_img,style_img,result_image], axis=0)  
vgg16_input
```

```
Out[6]: <tf.Tensor 'concat:0' shape=(3, 302, 403, 3) dtype=float32>
```

Import VGG16 model, and make a dictionary

```
In [7]: model = vgg16.VGG16(input_tensor=vgg16_input,
                             weights='imagenet', include_top=False)
model.summary()
#dictionary
model_dict = dict([(layer.name, layer.output) for layer in model.layers])
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(3, 302, 403, 3)	0
block1_conv1 (Conv2D)	(3, 302, 403, 64)	1792
block1_conv2 (Conv2D)	(3, 302, 403, 64)	36928
block1_pool (MaxPooling2D)	(3, 151, 201, 64)	0
block2_conv1 (Conv2D)	(3, 151, 201, 128)	73856
block2_conv2 (Conv2D)	(3, 151, 201, 128)	147584
block2_pool (MaxPooling2D)	(3, 75, 100, 128)	0
block3_conv1 (Conv2D)	(3, 75, 100, 256)	295168
block3_conv2 (Conv2D)	(3, 75, 100, 256)	590080
block3_conv3 (Conv2D)	(3, 75, 100, 256)	590080
block3_pool (MaxPooling2D)	(3, 37, 50, 256)	0
block4_conv1 (Conv2D)	(3, 37, 50, 512)	1180160
block4_conv2 (Conv2D)	(3, 37, 50, 512)	2359808
block4_conv3 (Conv2D)	(3, 37, 50, 512)	2359808
block4_pool (MaxPooling2D)	(3, 18, 25, 512)	0
block5_conv1 (Conv2D)	(3, 18, 25, 512)	2359808
block5_conv2 (Conv2D)	(3, 18, 25, 512)	2359808
block5_conv3 (Conv2D)	(3, 18, 25, 512)	2359808
block5_pool (MaxPooling2D)	(3, 9, 12, 512)	0
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

Check the dictionary

```
In [8]: model_dict
```

```
Out[8]: {'block1_conv1': <tf.Tensor 'block1_conv1/Relu:0' shape=(3, 302, 403, 64)
dtype=float32>,
'block1_conv2': <tf.Tensor 'block1_conv2/Relu:0' shape=(3, 302, 403, 64)
dtype=float32>,
'block1_pool': <tf.Tensor 'block1_pool/MaxPool:0' shape=(3, 151, 201, 6
4) dtype=float32>,
'block2_conv1': <tf.Tensor 'block2_conv1/Relu:0' shape=(3, 151, 201, 12
8) dtype=float32>,
'block2_conv2': <tf.Tensor 'block2_conv2/Relu:0' shape=(3, 151, 201, 12
8) dtype=float32>,
'block2_pool': <tf.Tensor 'block2_pool/MaxPool:0' shape=(3, 75, 100, 12
8) dtype=float32>,
'block3_conv1': <tf.Tensor 'block3_conv1/Relu:0' shape=(3, 75, 100, 256)
dtype=float32>,
'block3_conv2': <tf.Tensor 'block3_conv2/Relu:0' shape=(3, 75, 100, 256)
dtype=float32>,
'block3_conv3': <tf.Tensor 'block3_conv3/Relu:0' shape=(3, 75, 100, 256)
dtype=float32>,
'block3_pool': <tf.Tensor 'block3_pool/MaxPool:0' shape=(3, 37, 50, 256)
dtype=float32>,
'block4_conv1': <tf.Tensor 'block4_conv1/Relu:0' shape=(3, 37, 50, 512)
dtype=float32>,
'block4_conv2': <tf.Tensor 'block4_conv2/Relu:0' shape=(3, 37, 50, 512)
dtype=float32>,
'block4_conv3': <tf.Tensor 'block4_conv3/Relu:0' shape=(3, 37, 50, 512)
dtype=float32>,
'block4_pool': <tf.Tensor 'block4_pool/MaxPool:0' shape=(3, 18, 25, 512)
dtype=float32>,
'block5_conv1': <tf.Tensor 'block5_conv1/Relu:0' shape=(3, 18, 25, 512)
dtype=float32>,
'block5_conv2': <tf.Tensor 'block5_conv2/Relu:0' shape=(3, 18, 25, 512)
dtype=float32>,
'block5_conv3': <tf.Tensor 'block5_conv3/Relu:0' shape=(3, 18, 25, 512)
dtype=float32>,
'block5_pool': <tf.Tensor 'block5_pool/MaxPool:0' shape=(3, 9, 12, 512)
dtype=float32>,
'input_1': <tf.Tensor 'concat:0' shape=(3, 302, 403, 3) dtype=float32>}
```

Vgg16 is usually used to classificate, however, here we need to delete the classification purpose layers.

In this case, we decided to use 6 convolution layers to get the image's feature and one layers as output style feature

```
In [9]: # 6 convolution layers and one output style feature
feature_layers = ['block1_conv1', 'block1_conv2',
                  'block2_conv1', 'block3_conv1',
                  'block4_conv1', 'block5_conv1']
nfeature_layer = 6
output_layers = ['block5_conv2']
```

Using dictionary and layers we choosed late step to gerenate features

```
In [10]: features = model_dict['block5_conv2']
ori_features = features[0, :, :, :]
new_features = features[2, :, :, :]
```

We need a optimizer to minimize the loss between style image and new image. After reading the "Optimizers Guide" (<https://qiskit.org/documentation/aqua/optimizers.html>), we decided to use the Limited-memory Broyden-Fletcher-Goldfarb-Shanno Bound, because this optimizer has a direct callable function by scipy (https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.fmin_l_bfgs_b.html) (https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.fmin_l_bfgs_b.html)

```
In [11]: from scipy.optimize import fmin_l_bfgs_b
# x, f, d= fmin_l_bfgs_b(func, x0)

# x : array_like(Estimated position of the minimum.)
# f : float(Value of func at the minimum.)
# d : dict(Information dictionary.)
# func : callable f(x,*args) Function to minimise.
# x0 : ndarray Initial guess.
```

Here we decided to use the error square as loss and gradients descent which we learned in lecture.

Define Output Loss Function, using error square (keras.backend.square(x) <https://keras.io/backend/> (<https://keras.io/backend/>))

```
In [12]: # we has total 7 layers, the last ouput layer get smallest weight
output_weight=0.03
output_loss = output_weight * K.sum(K.square(new_features - ori_features))
```

Define Style Loss Function, using Gram Matrices

A Gram matrix of vectors a_1, \dots, a_n is a matrix G

s.t. $G = \langle a_i, a_j \rangle$ for all i, j if vectors a_1, \dots, a_n are columns of a matrix A , then $G = A^T A$ a Gram matrix is Positive Definite and Symmetric if vectors a_1, \dots, a_n are the rows of A (A would be so-called "Data Matrix"), then $G = A A^T$, and it's called left Gram matrix

Then still using error square to caclutate 2 matrices loss

```
In [13]: def style_loss_function(style_features, new_features):
#overturn style_features tensor
style_features = tensorflow.reshape(style_features,
                                     [-1, style_features.shape[-1]])

# G=AAT
style_gram_mtr = tensorflow.matmul(style_features, style_features,
                                   transpose_a=True, transpose_b=False)

new_features = tensorflow.reshape(new_features,
                                   [-1, style_features.shape[-1]])

# G=AAT
output_gram_mtr = tensorflow.matmul(new_features, new_features,
                                    transpose_a=True, transpose_b=False)

# Error square to caculate 2 matrices loss
style_loss = K.sum(K.square(style_gram_mtr - output_gram_mtr))
return style_loss
```

Define How to Caculate the initial gradients, using keras.backend.gradients(loss, variables)

<https://keras.io/backend/#gradients> (<https://keras.io/backend/#gradients>)

```
In [14]: for layer_name in feature_layers:
features = model_dict[layer_name]
style_features = features[1, :, :, :]
new_features = features[2, :, :, :]
style_loss = style_loss_function(style_features, new_features)
# 1 output layer, 6 convolution layers(features)
output_loss = output_loss + (1 / nfeature_layer) * style_loss
gradient = K.gradients(output_loss, result_image)
loss_list = [output_loss]
loss_list += gradient #Record gradient
keras_function = K.function([result_image], loss_list)
```

With loss function and initial gradients, now we define the callable function to minimise using in optimizer.

```
In [15]: class func(object):
def __init__(self):
self.loss_value = None

def loss(self, x):
x = x.reshape((1, ncol, nrow, 3))
result = keras_function([x])
self.loss_value = result[0]
self.grad_values = np.array(result[1:]).flatten().astype('float64')
return self.loss_value

def grad(self, x):
grad_values = np.copy(self.grad_values)
self.loss_value = None
self.grad_values = None
return grad_values
function = func()
```

Using original image as the input

```
In [16]: # run scipy-based optimization (L-BFGS) over the pixels of the generated image
# so as to minimize the neural style loss
X = load_img(ori_path, target_size=(ncol, nrow))
X = np.expand_dims(img_to_array(X), axis=0)
X = vgg16.preprocess_input(X)
```

Start iterating using Limited-memory Broyden-Fletcher-Goldfarb-Shanno Bound


```
In [17]: nitr = 400
neval = 20 #Maximum number of function evaluations.
loss = []
for i in range(nitr):
    print('Iteration = {}'.format(i))
    if i == 0:
        opt_loss = float('Inf') #initial min loss
    X, func_min, d = fmin_l_bfgs_b(function.loss, X.flatten(), fprime=function.grad)
    loss.append(func_min)
    print(d['grad'])
    if func_min < opt_loss:
        opt_loss = func_min
        opt_X = X.copy()
```

```
Iteration = 0
[ 1.12871153e+13  9.67454595e+12  3.44800965e+12 ...  1.10163237e+12
 -4.79169118e+11 -1.46461442e+11]
Iteration = 1
[-1.38585270e+12 -7.72080337e+11  1.25274489e+12 ...  2.99249140e+11
 2.53802136e+11  2.26965545e+10]
Iteration = 2
[ 1.42048231e+12  3.60088601e+11 -3.27851016e+11 ...  7.72686807e+10
 2.60116019e+09  9.72613386e+10]
Iteration = 3
[2.75204997e+11  6.76066263e+10  2.63078314e+11 ...  6.59419955e+10
 1.10414971e+11  8.19095142e+10]
Iteration = 4
[ 2.99091149e+10 -6.28181565e+10 -3.54626109e+10 ...  1.47346883e+10
 8.29418865e+10  7.99861637e+10]
Iteration = 5
[-7.14945331e+10  3.54930065e+11  4.05361787e+10 ...  1.99629906e+10
 2.46782198e+10  5.56122235e+10]
Iteration = 6
[-1.31225500e+11  1.55001001e+11  5.00000000e+10 ...  0.00000000e+00
 0.00000000e+00  0.00000000e+00]
```

Plot the minimum loss each iteration.


```
In [28]: plt.plot(loss)
plt.grid()
plt.xlabel('Iteration')
plt.ylabel('loss')
```

```
Out[28]: Text(0, 0.5, 'loss')
```



Then convert output array to image.

VGG-16 was trained using Caffe, and Caffe uses OpenCV to load images which uses BGR by default, so both VGG models are expecting BGR images.


The official mean that gets used in the vgg16.py file is [123.68, 116.779, 103.939]
(<https://forums.fast.ai/t/how-is-vgg16-mean-calculated/4577/9> (<https://forums.fast.ai/t/how-is-vgg16-mean-calculated/4577/9>))

So we need to convert BGR to RGB before imshow

```
In [19]: opt_X = opt_X.reshape((ncol, nrow, 3))
vgg_mean = [103.939, 116.779, 123.68]
for i in range(3):
    opt_X[:, :, i] = opt_X[:, :, i] + vgg_mean[i]
plt.imshow(opt_X)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] f or floats or [0..255] for integers).

Out[19]: <matplotlib.image.AxesImage at 0x7f022aa3c0b8>




Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
In [20]: opt_X = opt_X[:, :, ::-1]
opt_X = np.clip(opt_X, 0, 255)
plt.imshow(opt_X)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] f or floats or [0..255] for integers).

Out[20]: <matplotlib.image.AxesImage at 0x7f022aaafc518>




Convert the numpy arrays to uint8 before passing them to Image.fromarray

np array should have data type as uint8

(<https://stackoverflow.com/questions/10443295/combine-3-separate-numpy-arrays-to-an-rgb-image-in-python> (<https://stackoverflow.com/questions/10443295/combine-3-separate-numpy-arrays-to-an-rgb-image-in-python>))

```
In [24]: opt_X = opt_X.astype('uint8')  
plt.figure(figsize=(10,10))  
plt.imshow(opt_X)
```

Out[24]: <matplotlib.image.AxesImage at 0x7f022ab5b5c0>



```
In [22]: # Ori
plt.figure(figsize=(20,20))
plt.subplot(1,3,1)
plt.title("ori image",fontsize=20)
plt.axis('off')
plt.imshow(ori_image)
# style
plt.subplot(1,3,2)
plt.title("style",fontsize=20)
img_style = load_img(style_path)
plt.axis('off')
plt.imshow(style_image)
# new painting
plt.subplot(1,3,3)
plt.title("new painting",fontsize=20)
plt.axis('off')
plt.imshow(opt_X)
```

Out[22]: <matplotlib.image.AxesImage at 0x7f022ac37b70>



In []:

image-style-change (another example)

(This part we got inspiration and learning from <https://www.kaggle.com/basu369victor/style-transfer-deep-learning-algorithm> (<https://www.kaggle.com/basu369victor/style-transfer-deep-learning-algorithm>))

First, we have a original image and a style image . Then, we still use VGG16 but only pick the convolution layers for extracting image's features and delete the layers for classification. Finally, we will get a new picture combine the style of the style image and the objects from original image .

(This part we got inspiration and learning from <https://www.kaggle.com/basu369victor/style-transfer-deep-learning-algorithm> (<https://www.kaggle.com/basu369victor/style-transfer-deep-learning-algorithm>))

```
In [1]: import os
import numpy as np
import tensorflow
from tensorflow.keras import applications
from tensorflow.keras.applications import vgg16
from tensorflow.keras.models import Model
from scipy.optimize import fmin_l_bfgs_b
from tensorflow.keras.preprocessing.image import load_img, save_img, img_to_array
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

```
/anaconda3/lib/python3.6/site-packages/h5py/__init__.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
from ._conv import register_converters as _register_converters
```

```
In [2]: import tensorflow.keras.backend as K
K.clear_session()
ori_path = './Edouard_Manet_49.jpg'
style_path = './Joan_Miro_38.jpg'
```

Generation image's size

```
In [3]: ori_image = load_img(ori_path)
nrow, ncol = ori_image.size
nrow = nrow // 2
ncol = ncol // 2
print(str(nrow), str(ncol))
```

486 590

Show 2 images

```
In [4]: ori_image = load_img(ori_path,target_size=(ncol, nrow))
print(ori_image.size)
style_image = load_img(style_path, target_size=(ncol, nrow))
plt.axis('off')
plt.title('Edouard_Manet')
plt.imshow(ori_image)
plt.show()
plt.axis('off')
plt.title('Joan_Miro')
plt.imshow(style_image)
```


(486, 590)

Edouard_Manet



Out[4]: <matplotlib.image.AxesImage at 0x132812f550>

Joan_Miro



Pre-Process the images as the VGG16 input.Using tensor to represent the image,and bulid a new imgae as ouput

```
In [5]: def preprocess(image):
        img = img_to_array(image)
        img = vgg16.preprocess_input(img[np.newaxis,:])
        img = K.variable(img)
        return img
ori_img = preprocess(ori_image)
style_img = preprocess(style_image)
result_image = K.placeholder((1,ncol, nrow,3)) # 'channels_last' format()
result_image.shape
```

WARNING:tensorflow:From /usr/local/lib/python3.5/dist-packages/tensorflow/python/ops/resource_variable_ops.py:435: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

```
Out[5]: TensorShape([Dimension(1), Dimension(590), Dimension(486), Dimension(3)])
```

Then intergrating original image, style image and result image.

```
In [6]: vgg16_input = K.concatenate([ori_img,style_img,result_image], axis=0)
        vgg16_input
```

```
Out[6]: <tf.Tensor 'concat:0' shape=(3, 590, 486, 3) dtype=float32>
```

Import VGG16 model, and make a dictionary

```
In [7]: model = vgg16.VGG16(input_tensor=vgg16_input,
                             weights='imagenet', include_top=False)
model.summary()
#dictionary
model_dict = dict([(layer.name, layer.output) for layer in model.layers])
```

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(3, 590, 486, 3)	0
block1_conv1 (Conv2D)	(3, 590, 486, 64)	1792
block1_conv2 (Conv2D)	(3, 590, 486, 64)	36928
block1_pool (MaxPooling2D)	(3, 295, 243, 64)	0
block2_conv1 (Conv2D)	(3, 295, 243, 128)	73856
block2_conv2 (Conv2D)	(3, 295, 243, 128)	147584
block2_pool (MaxPooling2D)	(3, 147, 121, 128)	0
block3_conv1 (Conv2D)	(3, 147, 121, 256)	295168
block3_conv2 (Conv2D)	(3, 147, 121, 256)	590080
block3_conv3 (Conv2D)	(3, 147, 121, 256)	590080
block3_pool (MaxPooling2D)	(3, 73, 60, 256)	0
block4_conv1 (Conv2D)	(3, 73, 60, 512)	1180160
block4_conv2 (Conv2D)	(3, 73, 60, 512)	2359808
block4_conv3 (Conv2D)	(3, 73, 60, 512)	2359808
block4_pool (MaxPooling2D)	(3, 36, 30, 512)	0
block5_conv1 (Conv2D)	(3, 36, 30, 512)	2359808
block5_conv2 (Conv2D)	(3, 36, 30, 512)	2359808
block5_conv3 (Conv2D)	(3, 36, 30, 512)	2359808
block5_pool (MaxPooling2D)	(3, 18, 15, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		
=====		

Check the dictionary


```
In [8]: model_dict
```

```
Out[8]: {'block1_conv1': <tf.Tensor 'block1_conv1/Relu:0' shape=(3, 590, 486, 64)
dtype=float32>,
        'block1_conv2': <tf.Tensor 'block1_conv2/Relu:0' shape=(3, 590, 486, 64)
dtype=float32>,
        'block1_pool': <tf.Tensor 'block1_pool/MaxPool:0' shape=(3, 295, 243, 6
4) dtype=float32>,
        'block2_conv1': <tf.Tensor 'block2_conv1/Relu:0' shape=(3, 295, 243, 12
8) dtype=float32>,
        'block2_conv2': <tf.Tensor 'block2_conv2/Relu:0' shape=(3, 295, 243, 12
8) dtype=float32>,
        'block2_pool': <tf.Tensor 'block2_pool/MaxPool:0' shape=(3, 147, 121, 12
8) dtype=float32>,
        'block3_conv1': <tf.Tensor 'block3_conv1/Relu:0' shape=(3, 147, 121, 25
6) dtype=float32>,
        'block3_conv2': <tf.Tensor 'block3_conv2/Relu:0' shape=(3, 147, 121, 25
6) dtype=float32>,
        'block3_conv3': <tf.Tensor 'block3_conv3/Relu:0' shape=(3, 147, 121, 25
6) dtype=float32>,
        'block3_pool': <tf.Tensor 'block3_pool/MaxPool:0' shape=(3, 73, 60, 256)
dtype=float32>,
        'block4_conv1': <tf.Tensor 'block4_conv1/Relu:0' shape=(3, 73, 60, 512)
dtype=float32>,
        'block4_conv2': <tf.Tensor 'block4_conv2/Relu:0' shape=(3, 73, 60, 512)
dtype=float32>,
        'block4_conv3': <tf.Tensor 'block4_conv3/Relu:0' shape=(3, 73, 60, 512)
dtype=float32>,
        'block4_pool': <tf.Tensor 'block4_pool/MaxPool:0' shape=(3, 36, 30, 512)
dtype=float32>,
        'block5_conv1': <tf.Tensor 'block5_conv1/Relu:0' shape=(3, 36, 30, 512)
dtype=float32>,
        'block5_conv2': <tf.Tensor 'block5_conv2/Relu:0' shape=(3, 36, 30, 512)
dtype=float32>,
        'block5_conv3': <tf.Tensor 'block5_conv3/Relu:0' shape=(3, 36, 30, 512)
dtype=float32>,
        'block5_pool': <tf.Tensor 'block5_pool/MaxPool:0' shape=(3, 18, 15, 512)
dtype=float32>,
        'input_1': <tf.Tensor 'concat:0' shape=(3, 590, 486, 3) dtype=float32>}
```

Vgg16 is usually used to classificate, however, here we need to delete the classification purpose layers.

In this case, we decided to use 6 convolution layers to get the image's feature and one layers as output style feature

```
In [9]: # 6 convolution layers and one output style feature
feature_layers = ['block1_conv1', 'block1_conv2',
                  'block2_conv1', 'block3_conv1',
                  'block4_conv1', 'block5_conv1']
nfeature_layer = 6
output_layers = ['block5_conv2']
```

Using dictionary and layers we choosed late step to gerenate features

```
In [10]: features = model_dict['block5_conv2']
ori_features = features[0, :, :, :]
new_features = features[2, :, :, :]
```

We need a optimizer to minimize the loss between stype image and new image. After reading the "Optimizers Guide" (<https://qiskit.org/documentation/aqua/optimizers.html>), we decided to use the Limited-memory Broyden-Fletcher-Goldfarb-Shanno Bound, because this optimizer has a direct callable function by scipy (https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.fmin_l_bfgs_b.html) (https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.fmin_l_bfgs_b.html)

```
In [11]: from scipy.optimize import fmin_l_bfgs_b
# x, f, d= fmin_l_bfgs_b(func, x0)

# x : array_like(Estimated position of the minimum.)
# f : float(Value of func at the minimum.)
# d : dict(Information dictionary.)
# func : callable f(x,*args) Function to minimise.
# x0 : ndarray Initial guess.
```

Here we decided to use the error square as loss and gradients descent which we learned in lecture.

Define Output Loss Function, using error square (keras.backend.square(x) <https://keras.io/backend/>) (<https://keras.io/backend/>)

```
In [12]: # we has total 7 layers, the last ouput layer get smallest weight
output_weight=0.03
output_loss = output_weight * K.sum(K.square(new_features - ori_features))
```

Define Style Loss Function, using Gram Matrices

A Gram matrix of vectors a_1, \dots, a_n is a matrix G

s.t. $G = \langle a_i, a_j \rangle$ for all i, j if vectors a_1, \dots, a_n are columns of a matrix A , then $G = A^T A$ a Gram matrix is Positive Definite and Symmetric if vectors a_1, \dots, a_n are the rows of A (A would be so-called "Data Matrix"), then $G = A A^T$, and it's called left Gram matrix

Then still using error square to caclutate 2 matrices loss

```
In [13]: def style_loss_function(style_features, new_features):
#overturn style_features tensor
style_features = tensorflow.reshape(style_features,
                                     [-1, style_features.shape[-1]])

# G=AAT
style_gram_mtr = tensorflow.matmul(style_features, style_features,
                                   transpose_a=True, transpose_b=False)

new_features = tensorflow.reshape(new_features,
                                   [-1, style_features.shape[-1]])

# G=AAT
output_gram_mtr = tensorflow.matmul(new_features, new_features,
                                    transpose_a=True, transpose_b=False)

# Error square to caculate 2 matrices loss
style_loss = K.sum(K.square(style_gram_mtr - output_gram_mtr))
return style_loss
```

Define How to Caculate the initial gradients, using keras.backend.gradients(loss, variables)

<https://keras.io/backend/#gradients> (<https://keras.io/backend/#gradients>)

```
In [14]: for layer_name in feature_layers:
features = model_dict[layer_name]
style_features = features[1, :, :, :]
new_features = features[2, :, :, :]
style_loss = style_loss_function(style_features, new_features)
# 1 output layer, 6 convolution layers(features)
output_loss = output_loss + (1 / nfeature_layer) * style_loss
gradient = K.gradients(output_loss, result_image)
loss_list = [output_loss]
loss_list += gradient #Record gradient
keras_function = K.function([result_image], loss_list)
```

With loss function and initial gradients, now we define the callable function to minimise using in optimizer.

```
In [15]: class func(object):
def __init__(self):
self.loss_value = None

def loss(self, x):
x = x.reshape((1, ncol, nrow, 3))
result = keras_function([x])
self.loss_value = result[0]
self.grad_values = np.array(result[1:]).flatten().astype('float64')
return self.loss_value

def grad(self, x):
grad_values = np.copy(self.grad_values)
self.loss_value = None
self.grad_values = None
return grad_values
function = func()
```

Using original image as the input

```
In [16]: # run scipy-based optimization (L-BFGS) over the pixels of the generated image
# so as to minimize the neural style loss
X = load_img(ori_path, target_size=(ncol, nrow))
X = np.expand_dims(img_to_array(X), axis=0)
X = vgg16.preprocess_input(X)
```

Start iterating using Limited-memory Broyden-Fletcher-Goldfarb-Shanno Bound


```
In [17]: nitr = 400
neval = 20 #Maximum number of function evaluations.
loss = []
for i in range(nitr):
    print('Iteration = {}'.format(i))
    if i == 0:
        opt_loss = float('Inf') #initial min loss
    X, func_min, d = fmin_l_bfgs_b(function.loss, X.flatten(), fprime=function.loss)
    loss.append(func_min)
    print(d['grad'])
    if func_min < opt_loss:
        opt_loss = func_min
        opt_X = X.copy()
```

```
Iteration = 0
[-1.13819259e+14  2.11088499e+13  7.24103636e+13 ...  5.96972038e+12
 1.85995991e+13  1.44182818e+12]
Iteration = 1
[ 9.70059153e+12  1.30129005e+13  1.63229123e+10 ...  4.57716374e+12
-6.35074773e+12 -1.16962584e+12]
Iteration = 2
[-7.88490564e+13 -6.56955681e+13 -6.55839828e+13 ...  1.81990851e+12
-6.15783072e+12 -9.84203788e+11]
Iteration = 3
[-5.32359309e+12  2.57250443e+13  1.50716348e+13 ...  1.94227130e+12
 4.34667520e+09 -1.62685688e+12]
Iteration = 4
[-7.81255893e+13 -6.35482690e+13 -9.43444428e+13 ...  3.86211303e+12
-1.52068869e+12 -1.64111424e+12]
Iteration = 5
[-1.63219086e+13 -1.65134074e+11 -9.60866497e+12 ...  2.12884980e+12
-9.67161414e+11 -6.63994434e+11]
Iteration = 6
[-7.16000000e+13  1.11507515e+13  1.00000000e+13 ...  1.10571011e+12
```

Plot the minimum loss each iteration.

```
In [18]: plt.plot(loss)
plt.grid()
plt.xlabel('Iteration')
plt.ylabel('loss')
```

```
Out[18]: Text(0, 0.5, 'loss')
```



Then convert output array to image.

VGG-16 was trained using Caffe, and Caffe uses OpenCV to load images which uses BGR by default, so both VGG models are expecting BGR images.


The official mean that gets used in the vgg16.py file is [123.68, 116.779, 103.939]
(<https://forums.fast.ai/t/how-is-vgg16-mean-calculated/4577/9> (<https://forums.fast.ai/t/how-is-vgg16-mean-calculated/4577/9>))

So we need to covert BGR to RGB before imshow

```
In [19]: opt_X = opt_X.reshape((ncol, nrow, 3))
vgg_mean = [103.939, 116.779, 123.68]
for i in range(3):
    opt_X[:, :, i] = opt_X[:, :, i] + vgg_mean[i]
plt.imshow(opt_X)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] f or floats or [0..255] for integers).

Out[19]: <matplotlib.image.AxesImage at 0x7f3dd4526198>

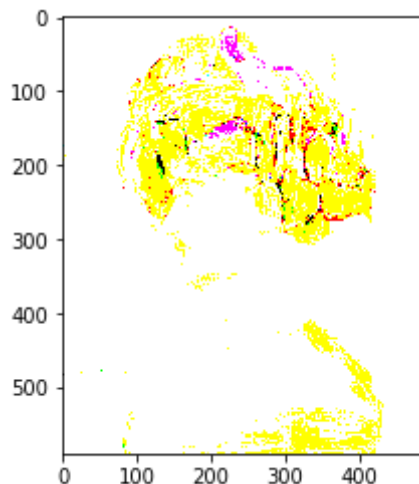


Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
In [20]: opt_X = opt_X[:, :, ::-1]
opt_X = np.clip(opt_X, 0, 255)
plt.imshow(opt_X)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] f or floats or [0..255] for integers).

Out[20]: <matplotlib.image.AxesImage at 0x7f3dd435b390>




Convert the numpy arrays to uint8 before passing them to Image.fromarray

np array should have data type as uint8

(<https://stackoverflow.com/questions/10443295/combine-3-separate-numpy-arrays-to-an-rgb-image-in-python> (<https://stackoverflow.com/questions/10443295/combine-3-separate-numpy-arrays-to-an-rgb-image-in-python>))


```
In [21]: opt_X = opt_X.astype('uint8')  
plt.figure(figsize=(10,10))  
plt.imshow(opt_X)
```

Out[21]: <matplotlib.image.AxesImage at 0x7f3dd4310a58>




```
In [22]: # Ori
plt.figure(figsize=(20,20))
plt.subplot(1,3,1)
plt.title("ori image",fontsize=20)
plt.axis('off')
plt.imshow(ori_image)
# style
plt.subplot(1,3,2)
plt.title("style",fontsize=20)
img_style = load_img(style_path)
plt.axis('off')
plt.imshow(style_image)
# new painting
plt.subplot(1,3,3)
plt.title("new painting",fontsize=20)
plt.axis('off')
plt.imshow(opt_X)
```

Out[22]: <matplotlib.image.AxesImage at 0x7f3dd42b91d0>



In []: